

```
[at3684@login-2-1 official]$ bash main.sh
20/12/02 22:41:05 INFO fs.TrashPolicyDefault: Moved: 'hdfs://dumbo/user/at3684/infrastructure/tax' to trash at: hdfs://dumbo/user/at3684/.Trash/Current
/user/at3684/infrastructure/tax1606966865027
rm: cannot remove `*.class': No such file or directory
Note: Clean.java uses or overrides a deprecated API.
Note: Recompile with -Xlint:deprecation for details.
added manifest
adding: Clean.class(in = 1377) (out= 797)(deflated 42%)
adding: CleanMapper.class(in = 5476) (out= 2785)(deflated 49%)
adding: CleanReducer.class(in = 1525) (out= 615)(deflated 59%)
20/12/02 22:41:18 WARN mapreduce.JobResourceUploader: Hadoop command-line option parsing not performed. Implement the Tool interface and execute your a
pplication with ToolRunner to remedy this.
20/12/02 22:41:19 INFO Input.FileInputFormat: Total input paths to process : 1
20/12/02 22:41:40 INFO hdfs.DFSCClient: Exception in createBlockOutputStream
java.io.IOException: Bad connect ack with firstBadLink as 10.0.255.213:50010
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.createBlockOutputStream(DFSOutputStream.java:1785)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.nextBlockOutputStream(DFSOutputStream.java:1683)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:790)
20/12/02 22:41:40 WARN hdfs.DFSCClient: Abandoning BP-1256044058-128.122.215.50-1440607644284:blk_1167383588_93862129
20/12/02 22:41:40 WARN hdfs.DFSCClient: Excluding datanode DatanodeInfoWithStorage[10.0.255.213:50010,DS-4222b439-c879-4567-8def-4b72087bc1ae,DISK]
20/12/02 22:42:00 INFO hdfs.DFSCClient: Exception in createBlockOutputStream
java.io.IOException: Bad connect ack with firstBadLink as 10.0.255.223:50010
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.createBlockOutputStream(DFSOutputStream.java:1785)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.nextBlockOutputStream(DFSOutputStream.java:1683)
    at org.apache.hadoop.hdfs.DFSOutputStream$DataStreamer.run(DFSOutputStream.java:790)
20/12/02 22:42:00 WARN hdfs.DFSCClient: Abandoning BP-1256044058-128.122.215.50-1440607644284:blk_1167383596_93862137
20/12/02 22:42:00 WARN hdfs.DFSCClient: Excluding datanode DatanodeInfoWithStorage[10.0.255.223:50010,DS-a4e6659e-8c9e-4493-802b-7da79a2ea6e4,DISK]
20/12/02 22:42:01 WARN hdfs.DFSCClient: Slow waitForAckedSegno took 40633ms (threshold=30000ms). File being written: /user/at3684/.staging/job_160434639
2376_7784/job.splitmetainfo, block: BP-1256044058-128.122.215.50-1440607644284:blk_1167383603_93862144, Write pipeline datanodes: null
20/12/02 22:42:01 INFO mapreduce.JobSubmitter: number of splits:1
20/12/02 22:42:04 INFO mapreduce.JobSubmitter: Submitting tokens for job: job_1604346392376_7784
20/12/02 22:42:04 INFO impl.YarnClientImpl: Submitted application application_1604346392376_7784
20/12/02 22:42:04 INFO mapreduce.Job: The url to track the job: http://babar.es.its.nyu.edu:8088/proxy/application_1604346392376_7784/
20/12/02 22:42:04 INFO mapreduce.Job: Running job: job_1604346392376_7784
20/12/02 22:42:08 INFO mapreduce.Job: Job job_1604346392376_7784 running in uber mode : false
20/12/02 22:42:08 INFO mapreduce.Job: map 0% reduce 0%
20/12/02 22:42:13 INFO mapreduce.Job: map 100% reduce 0%
20/12/02 22:42:18 INFO mapreduce.Job: map 100% reduce 100%
20/12/02 22:42:20 INFO mapreduce.Job: Job job_1604346392376_7784 completed successfully
20/12/02 22:42:20 INFO mapreduce.Job: Counters: 49
File System Counters
    FILE: Number of bytes read=79400
    FILE: Number of bytes written=467217
    FILE: Number of read operations=0
    FILE: Number of large read operations=0
    FILE: Number of write operations=0
    HDFS: Number of bytes read=4701295
    HDFS: Number of bytes written=151618
    HDFS: Number of read operations=6
    HDFS: Number of large read operations=0
    HDFS: Number of write operations=2
Job Counters
    Launched map tasks=1
    Launched reduce tasks=1
    Data-local map tasks=1
    Total time spent by all maps in occupied slots (ms)=10004
    Total time spent by all reduces in occupied slots (ms)=15228
    Total time spent by all map tasks (ms)=2501
    Total time spent by all reduce tasks (ms)=2538
    Total vcore-milliseconds taken by all map tasks=2501
    Total vcore-milliseconds taken by all reduce tasks=2538
    Total megabyte-milliseconds taken by all map tasks=10244096
    Total megabyte-milliseconds taken by all reduce tasks=15593472
Map-Reduce Framework
    Map input records=3191
    Map output records=3141
    Map output bytes=164182
    Map output materialized bytes=79396
    Input split bytes=145
    Combine input records=0
    Combine output records=0
    Reduce input groups=3141
    Reduce shuffle bytes=79396
    Reduce input records=3141
    Reduce output records=3141
    Spilled Records=6282
    Shuffled Maps =1
    Failed Shuffles=0
    Merged Map outputs=1
    GC time elapsed (ms)=374
    CPU time spent (ms)=8410
    Physical memory (bytes) snapshot=1488326656
```